

Image Classification for Cervical Spine CT Scans **Final Presentation**

Aniket Lachyankar & Satwik Kamarthi
CS7150 - Professor David Bau
December 5, 2022

Table of **C**ontents

01

Introduction

Our motivation of project as well as project goals

02

Data

Overview of dataset and preprocessing steps

03


Model Selection

Models we chose to research and implement

04

Results & Future Work

Performance analysis and possible next steps to improve performance.






01

Introduction


Our **motivation**

- ❖ An Image is Worth 16x16 Words: Transformers For Image Recognition At Scale
- ❖ Cervical Spinal Fractures domain
- ❖ Quick detection is essential to prevent neurological deterioration and paralysis after trauma



Project Goal

Understand the performance of top models in the medical imaging domain by comparing a ResNet model and Vision Transformer model and evaluate how well they can handle a CT scan classification task.



The background is a dark blue gradient. On the left, there is a large, glowing orange and yellow shape. Below it is a smaller, glowing red and purple sphere. On the right, there is a large, glowing purple and red shape with a smaller, glowing red and purple sphere inside it. In the center, the text "02 Data" is displayed in white.


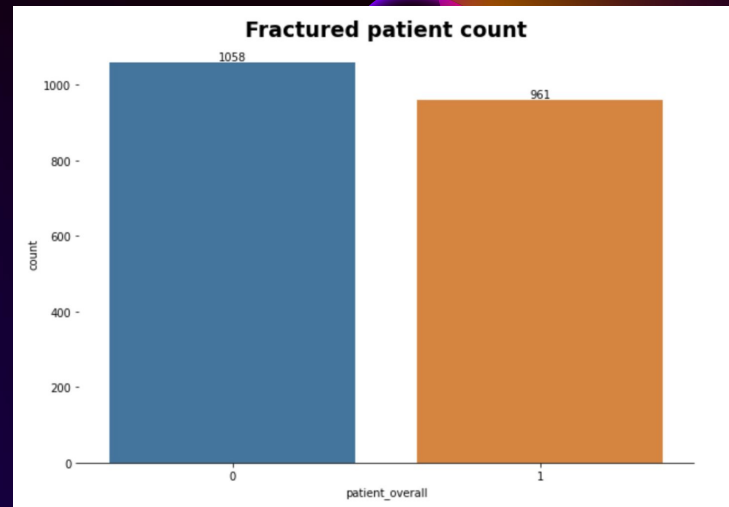
02 Data

Data overview

- ❖ The dataset we are using comes from RSNA 2022 Cervical Spine Fracture Detection Kaggle competition.
- ❖ Dataset comprised of Cervical Spine CT scans taken from 12 sites around the world by the Radiological Society of North America (RSNA)


- ❖ Training Data

- 2000 patients
 - 300-600 digital imaging files (.dcm) representing axial scans of fractures
- Distribution of fractures
- Fractured Vertebrae count



Data Preprocessing

- ❖ After meeting with a medical professional, we were advised that a better view to detect spinal fractures with our model may be the sagittal view (from the side).
- ❖ Since original dataset has all the images in axial view (top-down), we needed to convert all the axial images to sagittal images.
- ❖ Using methods adapted from pydicom library documentation:
 - Accessing DICOM image metadata
 - Create 3D array filled with original images from file
 - Manipulate & transform array to get the sagittal aspect and view from 3D array
- ❖ From 300-600 axial images we generated 25 sagittal images per patient in training set, and 100 sagittal images per patient in test set.
 - Total ~16K Train And Validation Images
 - Sample mainly from the middle of each scan as that is where the majority of each vertebrae can be seen.





03

Model Selection

ResNet

- ❖ Huggingface implementation
- ❖ Microsoft pre-trained weight checkpoint
- ❖ ResNet-50
 - Balance between performance and training time

Layer (type:depth-idx)	Output Shape	Param #
ResNetForImageClassification	[1, 1000]	--
└ResNetModel: 1-1	[1, 2048, 1, 1]	--
└ResNetEmbeddings: 2-1	[1, 64, 56, 56]	--
└ResNetConvLayer: 3-1	[1, 64, 112, 112]	9,536
└ResNetEncoder: 2-2	--	(recursive)
└ModuleList: 3-4	--	(recursive)
└ResNetEmbeddings: 2-3	--	(recursive)
└MaxPool2d: 3-3	[1, 64, 56, 56]	--
└ResNetEncoder: 2-4	[1, 2048, 7, 7]	--
└ModuleList: 3-4	--	(recursive)
└AdaptiveAvgPool2d: 2-5	[1, 2048, 1, 1]	--
└Sequential: 1-2	[1, 1000]	--
└Flatten: 2-6	[1, 2048]	--
└Linear: 2-7	[1, 1000]	2,049,000

Total params: 25,557,032
Trainable params: 25,557,032
Non-trainable params: 0
Total mult-adds (G): 4.09

Input size (MB): 0.60
Forward/backward pass size (MB): 177.83
Params size (MB): 102.23
Estimated Total Size (MB): 280.66

ResNet cont.

- ❖ Model hyperparameters inspired by PyTorch blogpost
- ❖ Small batch size
 - GPU memory size
- ❖ Ability to increase with given compute resources
- ❖ Because we have different number of labels, we use BCELoss for multilabel output
 - Use Sigmoid on each output class rather than softmax for all classes

Hyperparameter/Optimizer	Value
Loss Function	Binary Cross Entropy
Learning Rate	0.5
Learning Rate Scduler	Cosine Annealing
Learning Rate Warmup Epochs	5
Learning Rate Warmup Method	linear
Learning Rate Warmup Decay	0.01
Batch Size	24
Optimizer	SGD (Stochastic Gradient Descent)
SGD Momentum	0.9
Weight Decay	2e-05
Epochs	100

Vision Transformers

- ❖ Huggingface implementation
- ❖ Google and Huggingface pre-trained weight checkpoint
 - Avoid vanilla model
 - Trained on ImageNet-21k
- ❖ Stages and parameter sizes of model for 1 input image on right

Layer (type:depth-idx)	Output Shape	Param #
ViTForImageClassification	[1, 1000]	--
└ViTModel: 1-1	[1, 197, 768]	--
└ViTEmbeddings: 2-1	[1, 197, 768]	152,064
└ViTPatchEmbeddings: 3-1	[1, 196, 768]	590,592
└Dropout: 3-2	[1, 197, 768]	--
└ViTEncoder: 2-2	[1, 197, 768]	--
└ModuleList: 3-3	--	85,054,464
└LayerNorm: 2-3	[1, 197, 768]	1,536
└Linear: 1-2	[1, 1000]	769,000
Total params: 86,567,656		
Trainable params: 86,567,656		
Non-trainable params: 0		
Total mult-adds (M): 201.58		
Input size (MB): 0.60		
Forward/backward pass size (MB): 162.19		
Params size (MB): 345.66		
Estimated Total Size (MB): 508.46		

Vision Transformers cont.

- ❖ Adjust input images to (3,384, 384)
 - No models for images of size 512
- ❖ Similar to ResNet, we use BCELoss for calculating gradient and Sigmoid activation on output layer.


Hyperparameter/Optimizer	Value
Loss Function	Binary Cross Entropy
Learning Rate	3e-05
Batch Size	12
Optimizer	Adam
Adam Betas	(0.9, 0.999)
Adam Epsilon	1e-08
Epochs	10
Learning Rate Scduler	Linear




04

Results & Future Work


ViT Results



- ❖ We see that the model converges on training data after approximately 10 epochs.
 - Jagged validation loss suggests that the model may be overfitting.



ResNet Results







- ❖ Similar convergence after 10 epochs.
 - Seems like less overfitting based on validation loss.

Results Contd.

- ❖ Next, we tested the best trained model (lowest validation loss) on a test holdout set
 - 301 Sagittal Samples, all labeled with a C1 fracture.
- ❖ ResNet did not predict and correct samples.
 - Either dataset was not large enough, or we overfit the model on the training dataset
- ❖ ViT predicted 5 correct predictions out of 301 samples in the test set:
 - Low performance suggests overfitting in this case as well based on the loss graphs we see previously.

Results Contd.

- ❖ Below are the test samples that ViT correctly predicted with the C1 fracture label.



Analysis

A few takeaways from this project:

- Dataset selection/preprocessing is EXTREMELY important
 - Choices made regarding selecting data for training were the bulk of the project.
 - In our case, the used dataset size may have been too small for the complexity of the task.
 - A key task of our data preprocessing was selecting slices from the scan, which we did arbitrarily. Because of this and the nature of the labels, we cannot guarantee that
- Pre-Trained models may not contain trained parameters that work best for a domain vastly different from the pretraining.
 - Repurposed models that are usually for single label prediction didnt necessarily work well for complex multilabel prediction.

Future Work

- Improve input data
 - One major problem with the dataset provided is that we have no way of guaranteeing whether all the slices in a scan contain the fracture that is in the label.
 - Dataset came with bounding boxes for axial slices, maybe possible to use these for sagittal slices.
- Utilize models that can handle sequential data.
 - If we can pass all the slices of the scan as part of a single sequence of inputs, the model may have a better chance at analyzing multiple slices .



QUESTIONS?